

# REPORT OF IMPLEMENTATION

## Learning Algorithm

In this project was used Actor-Critic method, through the DDPG algorithm. DDPG means Deep Deterministic Policy Gradient and extend policy-based reinforcement learning methods to complex problems using deep neural networks.

DDPG is very similar to DQN method, using ReplayBuffer. A big difference is that DQN is used to discrete action spaces and DDPG is used to continuous action spaces.

In algorithms like DDPG that implement actor-critic method we have two neural networks, one is the actor and another is the critic.

Also, was used Experience Replay and Fixed Q-Targets methods to improve the agent's performance.

OUNoise Class is used to add noise to actions to promote exploration. It uses the Ornstein-Uhlenback process to achieve this.

## CODE

1- model.py

In this file was created the actor and critic classes, each one with your own neural network. The neural network was created using PyTorch.

2- ddpg\_agent.py

In this file was created the Agent class.

Looking for learning and improving the results, the agent implements the Experience Replay and Fixed Q-Targets methods.

Was created the ReplayBuffer class to implement the Experience Replay method

4- Training the agent

The agent was trained in 2.000 episodes and epsilon decay = 0.995

The environment is considered solved when the agent achieve the mean score  $\geq 13.0$ , but the train continues until finish the 2.000 episodes.

## ATTEMPT 1

### Model.py

Actor model architecture:

- three fully connected layers, receiving the state as input and actions as output.
- 33(Input-state\_size) x 400(hidden\_layer) x 300(hidden\_layer) x 4(output - action\_size)
- Was applied ReLu activation at hidden layers and Tanh activation at output layer.

Critic model architecture:


- three fully connected layers, receiving the state as input and return as output only one unit, the Q-values.
- 33(Input-state\_size) x 404(hidden\_layer+action\_size) x 300(hidden\_layer) x 1(output - Q-values)
- Was applied ReLu activation at hidden layers.

Hyperparameters used to train the agent:

```
BUFFER_SIZE = int(1e5) # replay buffer size
BATCH_SIZE = 128      # minibatch size
```

GAMMA = 0.99           # discount factor  
 TAU = 1e-3            # for soft update of target parameters  
 LR\_ACTOR = 1e-4       # learning rate of the actor  
 LR\_CRITIC = 1e-3      # learning rate of the critic  
 WEIGHT\_DECAY = 0      # L2 weight decay  
 UPDATE\_EVERY = 1      # how often to update the network  
 N\_EPISODES = 1000     # number of episodes

The result of the first attempt was really bad. The max scored achieved was around 2.5



## ATTEMPT 2

To second attempt I did some changes:

- Neural network architecture of the Actor
- Neural network architecture of the Critic
- Hyperparameters
- UoNoise

### Model.py

Actor model architecture:

- three fully connected layers, receiving the state as input and actions as output.
- 33(Input-state\_size) x 256(hidden\_layer) x 128(hidden\_layer) x 4(output - action\_size)
- Was applied ReLu activation at hidden layers and Tanh activation at output layer.

Critic model architecture:

- four fully connected layers, receiving the state as input and return as output only one unit, the Q-values.
- 33(Input-state\_size) x 132(hidden\_layer+action\_size) x 64(hidden\_layer) x 32(hidden\_layer) x 1(output - Q-values)

- Was applied ReLu activation at hidden layers.

Hyperparameters used to train the agent (Only modified hyperparameters):

```
BUFFER_SIZE = int(1e6) # replay buffer size
BATCH_SIZE = 1024      # minibatch size
```

"""Ornstein-Uhlenbeck process."""

Was necessary to change the noise process. Below the old and new version of code.

```
#dx = self.theta * (self.mu - x) + self.sigma * np.array([random.random() for i in range(len(x))])
dx = self.theta * (self.mu - x) + self.sigma * np.random.standard_normal(self.size)
```


## RESULT

With those architecture and hyper parameters was possible achieve the Average Score of the 30.19 in only 130 episodes.

Below the progress of the agent's learning.

Episode 10	Average Score: 1.19	Episode Score: 1.96
Episode 20	Average Score: 3.87	Episode Score: 10.60
Episode 30	Average Score: 7.39	Episode Score: 14.82
Episode 40	Average Score: 10.04	Episode Score: 19.07
Episode 50	Average Score: 12.07	Episode Score: 21.24
Episode 60	Average Score: 13.84	Episode Score: 24.01
Episode 70	Average Score: 15.72	Episode Score: 28.91
Episode 80	Average Score: 17.60	Episode Score: 30.17
Episode 90	Average Score: 19.45	Episode Score: 35.95
Episode 100	Average Score: 21.16	Episode Score: 36.73
Episode 110	Average Score: 24.76	Episode Score: 37.66
Episode 120	Average Score: 27.87	Episode Score: 37.90
Episode 130	Average Score: 30.19	Episode Score: 37.76

Environment solved in 130 episodes



### **Future ideas for improving agent's performance**

Looking for improving the agent's performance we suggest to try the following actions:

- Modifying the hyper parameters, looking for speed up training or increase the final score.
- Implement a different algorithm like PPO or D4PG.
- Modifying the model architecture by changing the number of layers or neurons.
- Change the update frequency of the networks in the step function.